

Data Deduplication with FLIC Manifests

NDNComm 2025

Marc Mosko

SupraLiminal Technologies

<https://github.com/mmosko/ccnx-dedup>


<https://github.com/mmosko/ccnpy>

<https://datatracker.ietf.org/doc/draft-irtf-icnrg-flic/07>

Outline

- Data Deduplication
- FLIC Manifests
- Methodology
- Results
- Conclusion

Data Deduplication [1]




Chunking splits the data at breakpoints.

Fingerprinting uses a fast hash to screen chunks.

A strong hash or byte comparison fully matches chunks.

[1] Karp, R.M. and Rabin, M.O., 1987. Efficient randomized pattern-matching algorithms. IBM journal of research and development, 31(2), pp.249-260.

Deduplication Matching



Karp-Rabin fingerprinting accommodates chunks moving

FastCDC [2] and Gear [3]

Algorithm 2. FastCDC8KB (with NC)

Input: data buffer, *src*; buffer length, *n*

Output: chunking breakpoint *i*

```

MaskS ← 0x0000d9f003530000LL;           // 15 '1' bits;
MaskA ← 0x0000d93003530000LL;           // 13 '1' bits;
MaskL ← 0x0000d90003530000LL;           // 11 '1' bits;
MinSize ← 2 KB; MaxSize ← 64 KB;
fp ← 0; i ← MinSize; NormalSize ← 8 KB;
if n ≤ MinSize then
    return n;
if n ≥ MaxSize then
    n ← MaxSize;
else if n ≤ NormalSize then
    NormalSize ← n;
for i < NormalSize; i++; do
    fp = (fp << 1) + Gear[ src[i] ];
    if ! ( fp & MaskS ) then
        return i;                //if the masked bits are all '0';
for i < n; i++; do
    fp = (fp << 1) + Gear[ src[i] ];
    if ! ( fp & MaskL ) then
        return i;                //if the masked bits are all '0';
return i;

```

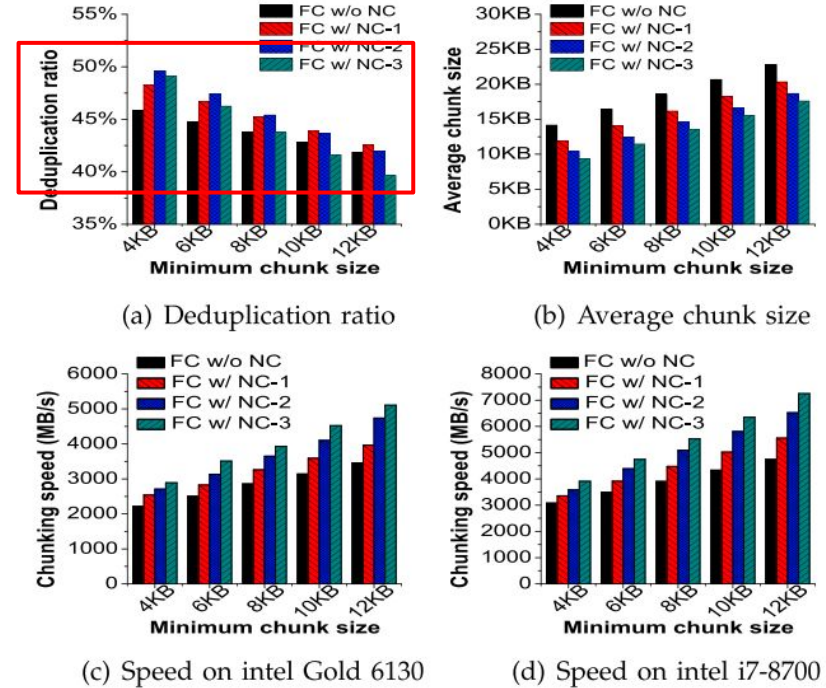


Fig. 13. Evaluation of comprehensive performance of normalized chunking with different normalization levels.

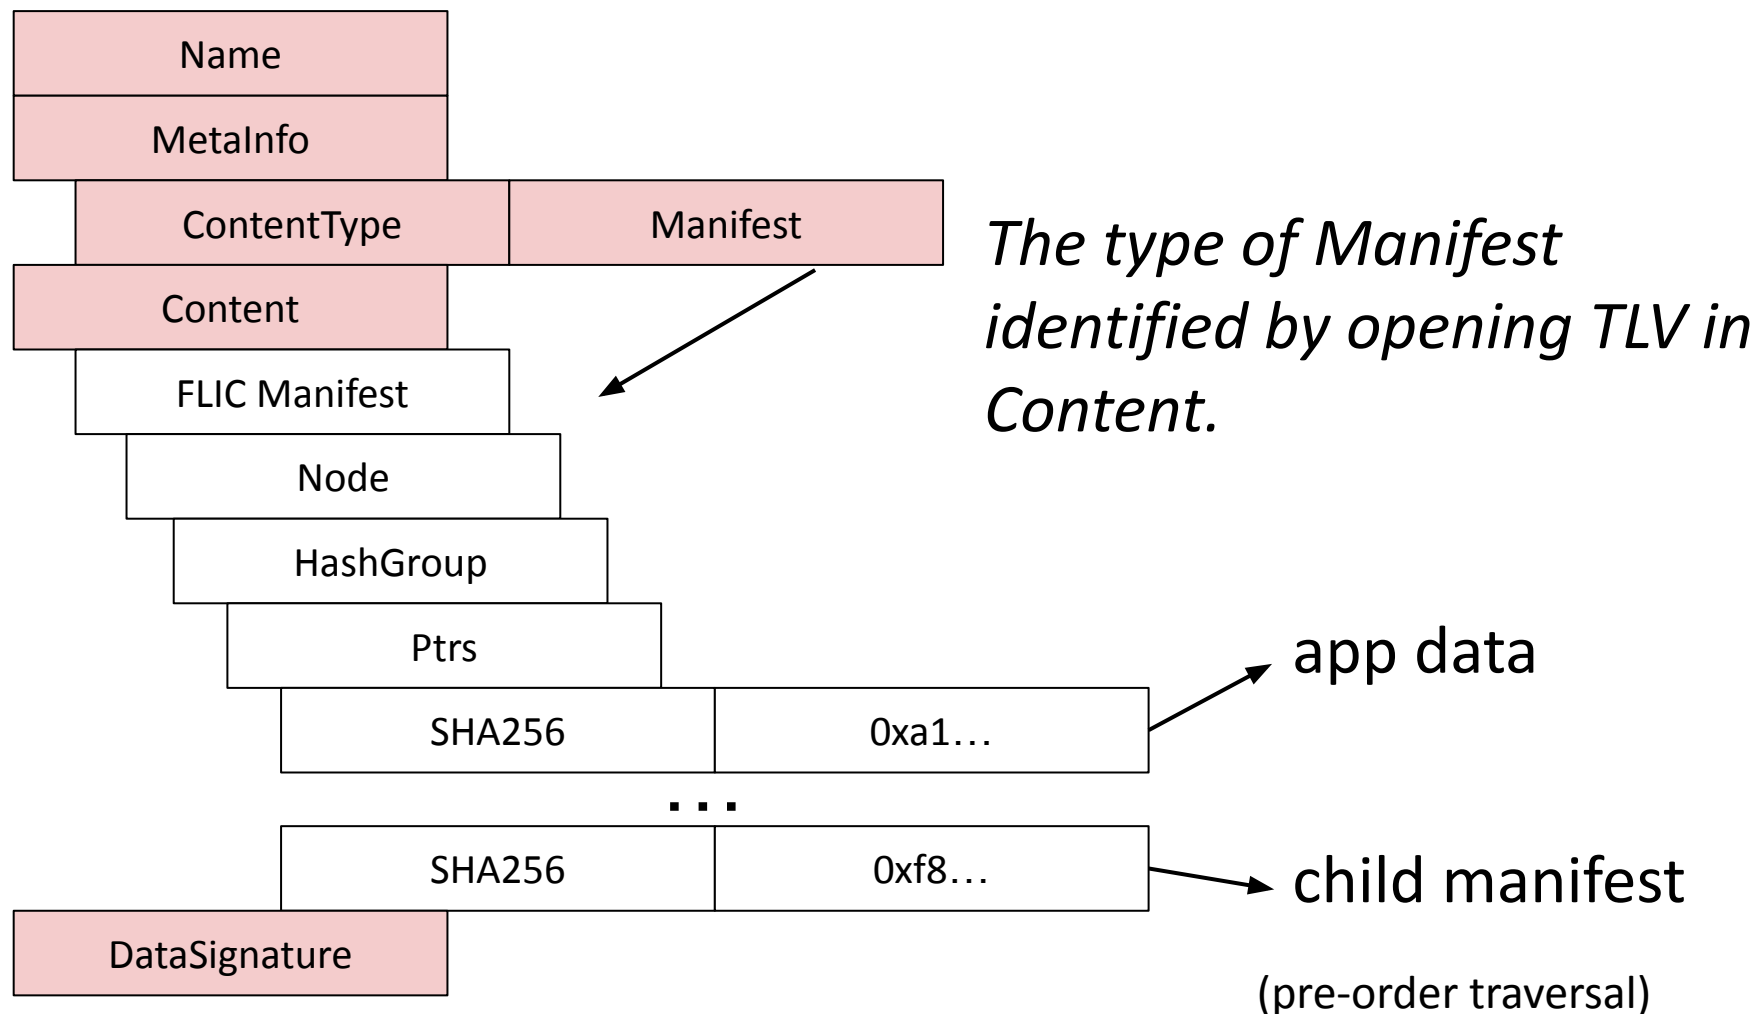
[2] Xia, Wen, Xiangyu Zou, Hong Jiang, Yukun Zhou, Chuanyi Liu, Dan Feng, Yu Hua, Yuchong Hu, and Yucheng Zhang. "The design of fast content-defined chunking for data deduplication based storage systems." IEEE Transactions on Parallel and Distributed Systems 31, no. 9 (2020): 2017-2031.

[3] Xia, W., Jiang, H., Feng, D., Tian, L., Fu, M. and Zhou, Y., 2014. Ddelta: A deduplication-inspired fast delta compression approach. Performance Evaluation, 79, pp.258-272.

FLIC Manifests

- File-Like Information Centric Manifests
 - <https://datatracker.ietf.org/doc/draft-irtf-icnrg-flic/07>
- Encodes a single object
- Reconstructed by concatenating Data object payloads bytes via pre-order traversal
- Follow-up work
 - Archives manifests of multiple FLICs
 - Object compression
 - Differential manifest encoding

FLIC Manifest Structure



Hashes To Names

- *Name Constructor* informs an Interest.
 - Segmented: Use a prefix plus a **T-V number**.
 - Prefix: A common prefix, use implicit hash.
 - Hash: CCNx Nameless Objects.
 - *Locators*, which become NDN routing hints.
 - Other Interest flags.
- A *HashGroup* specifies which NC to use.
- NCs are inherited.

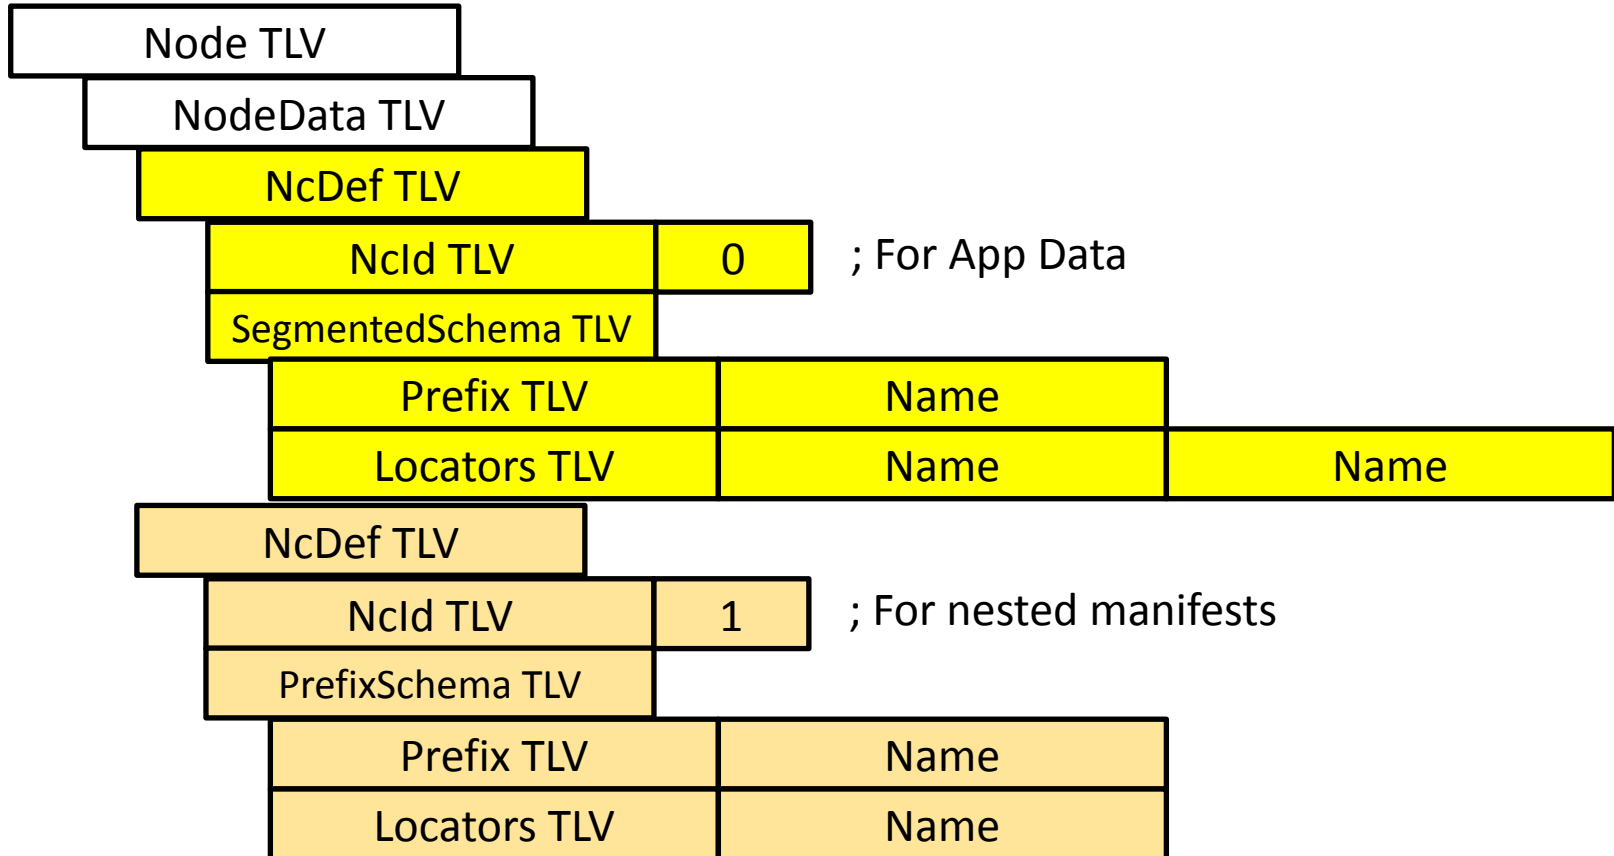
NDN Objectstore Naming

- NDN prefers that each Data object have a unique name (not counting the implicit hash).

common prefix + number (+hash) [Segmented]
common prefix (+hash) [Prefix]

- If the manifest has one prefix and the application data has a second prefix, then one must have two (or more) HashGroups.
- Manifests use a unique *ManifestId* name segment not a Segment Number.
- One must traverse the manifest in pointer order.

Unencrypted Manifest With NodeData



1. /foo/0x82=10, hash=0x0001
2. /foo/0x82=20, hash=0x0002
3. /foo/0x82=12, hash=0x0003
4. /bar/0x32=0, hash=0x0004
5. /bar/0x32=1, hash=0x0005
6. /bar/0x32=2, hash=0x0006

; manifests using ManifestId

; segmented app data

Manifest

Node

NodeData

NcDef NcId 1 SegmentedSchema Name '/foo' SuffixComponentType 0x82

NcDef NcId 2 SegmentedSchema Name '/bar' SuffixComponentType 0x32

HashGroup

GroupData NcId 1 StartSegmentId 10

AnnotatedPtrs

PointerBlock

Ptr HashValue(0x0001)

Ptr HashValue(0x0002) SegmentIdAnnotation(20)

Ptr HashValue(0x0003)

HashGroup

GroupData NcId 2 StartSegmentId 0

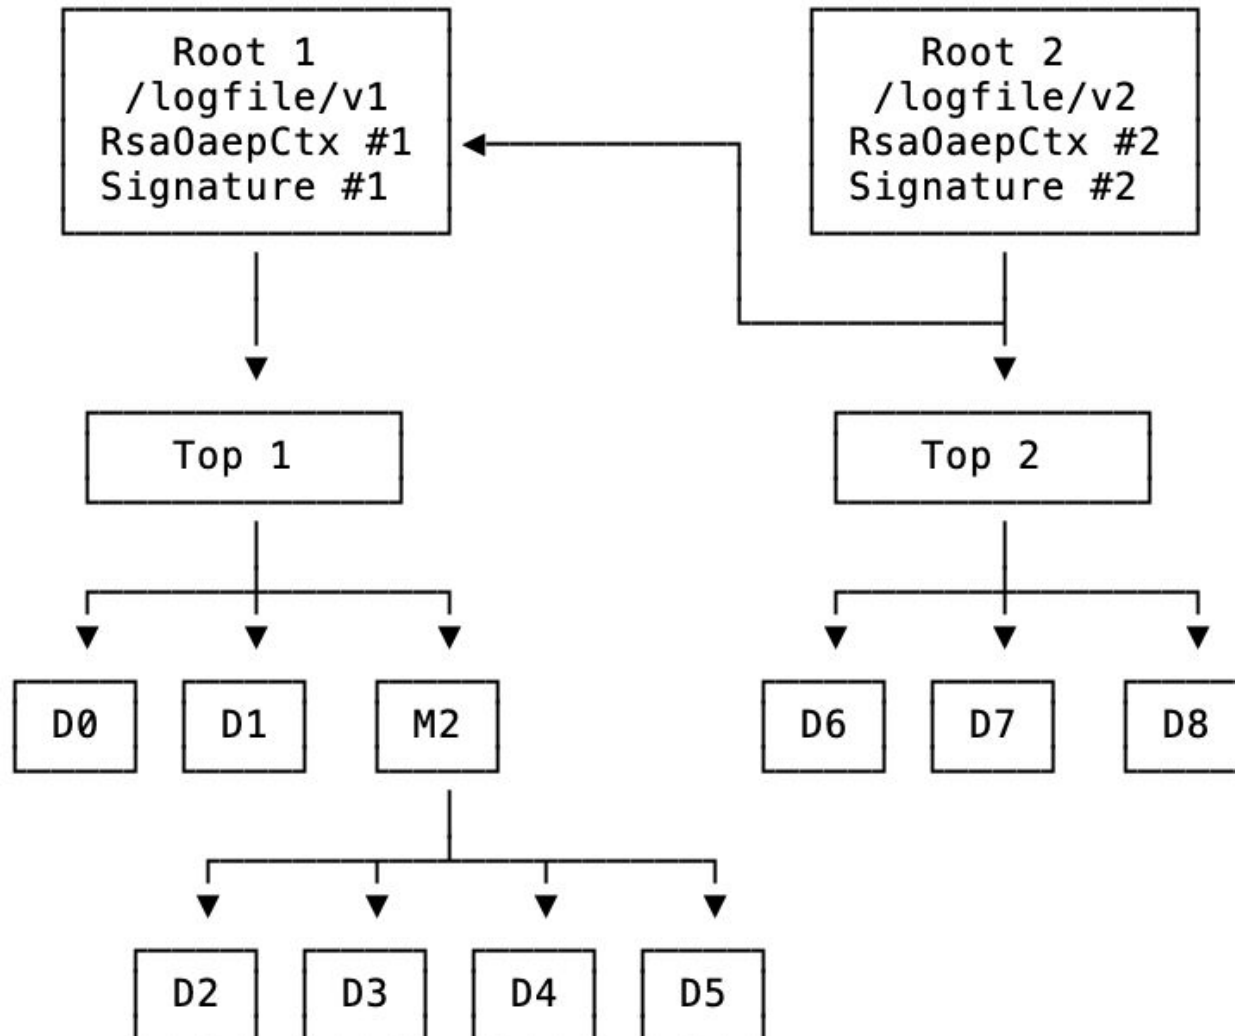
Ptrs

HashValue(0x0004)

HashValue(0x0005)

HashValue(0x0006)

Typical Manifest Structures



Other FLIC Features

- Encryption of Manifest
 - In-place encryption using AES
 - RSA-OAEP key wrapping with salt
 - KDF
- Metadata
 - Subtree size, Subtree digest
 - Annotated pointers
- Algorithm for one-pass bottom-up encoding.
- Growable, re-namable, re-signable.

Experimental Deduplication Results

- We applied Fastcdc data deduplication algorithm to a sampling of GNU source code distributions.
 - bison, emacs, and patch
 - (binutils and gcc show similar results)

Methodology

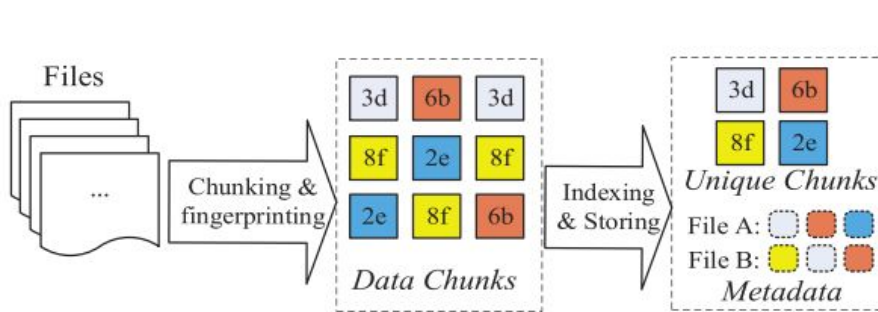


Fig. 1. General workflow of chunk-level data deduplication.[Xia et al.]




Fig. 2. The sliding window technique for the CDC algorithm. The hash value of the sliding window, fp , is computed via the Rabin algorithm (this is the *hashing stage* of CDC). If the lowest $\log_2 D$ bits of the hash value matches a threshold value r , i.e., $fp \bmod D = r$, this offset (i.e., the current position) is marked as a chunk cut-point (this is the *hash-judging stage* of CDC). [Xia et al.]

Each dedup chunk is about 2KB - 40KB.

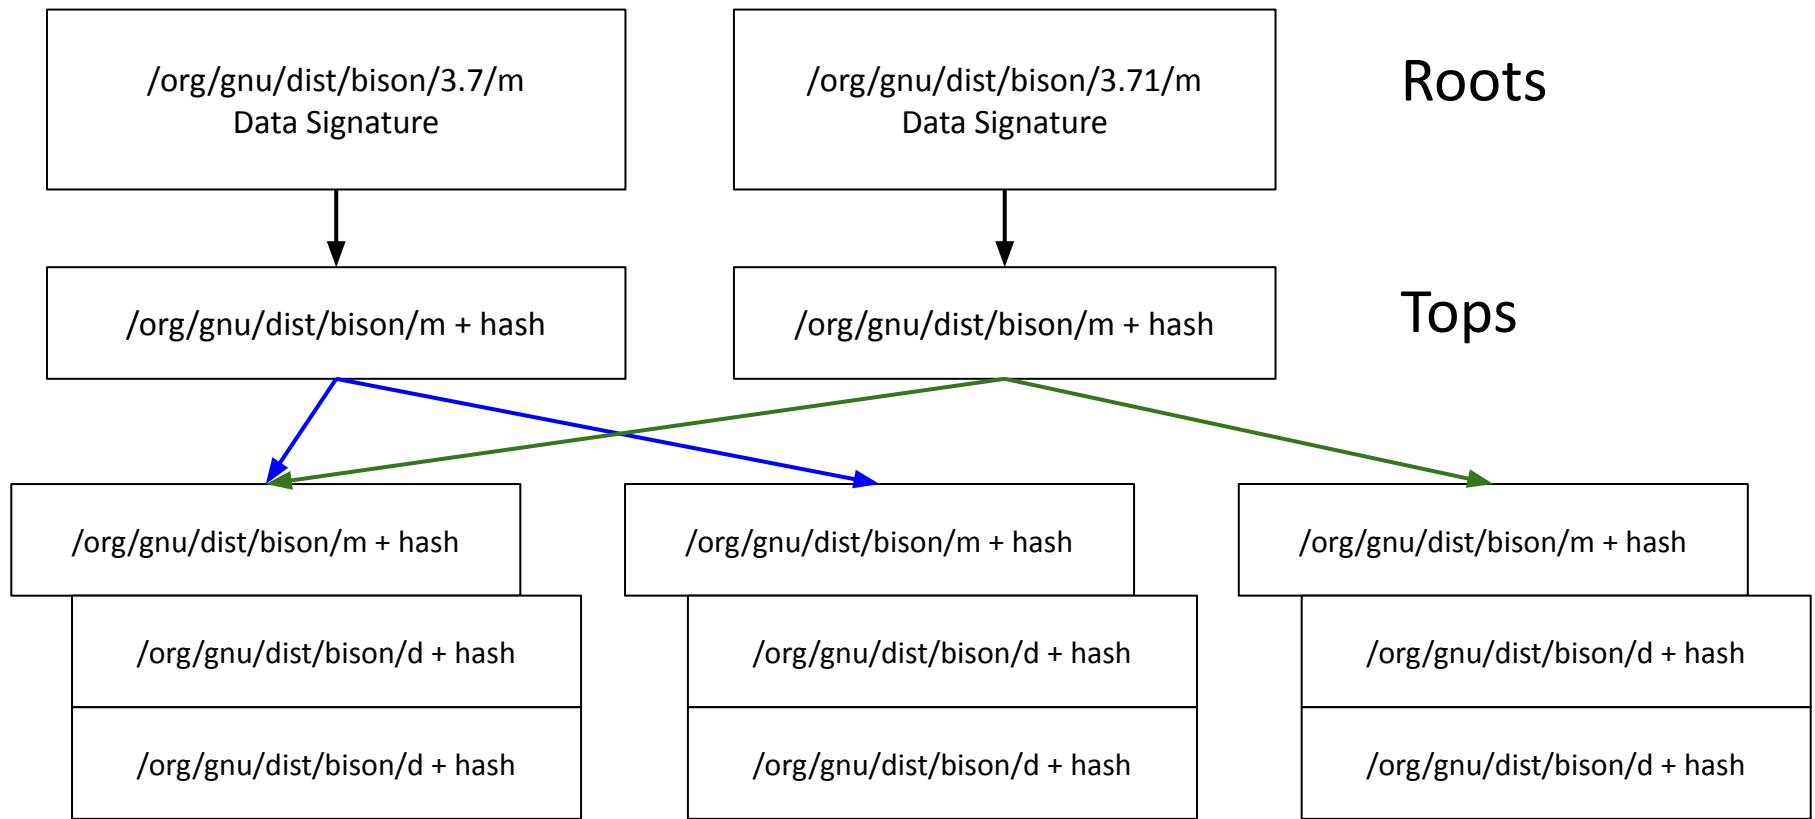
We encode each chunk as it's own mini-FLIC manifest.

Usually 1 FLIC manifest plus 7-30 data objects.

Each *tar* file has its own signed root manifest.


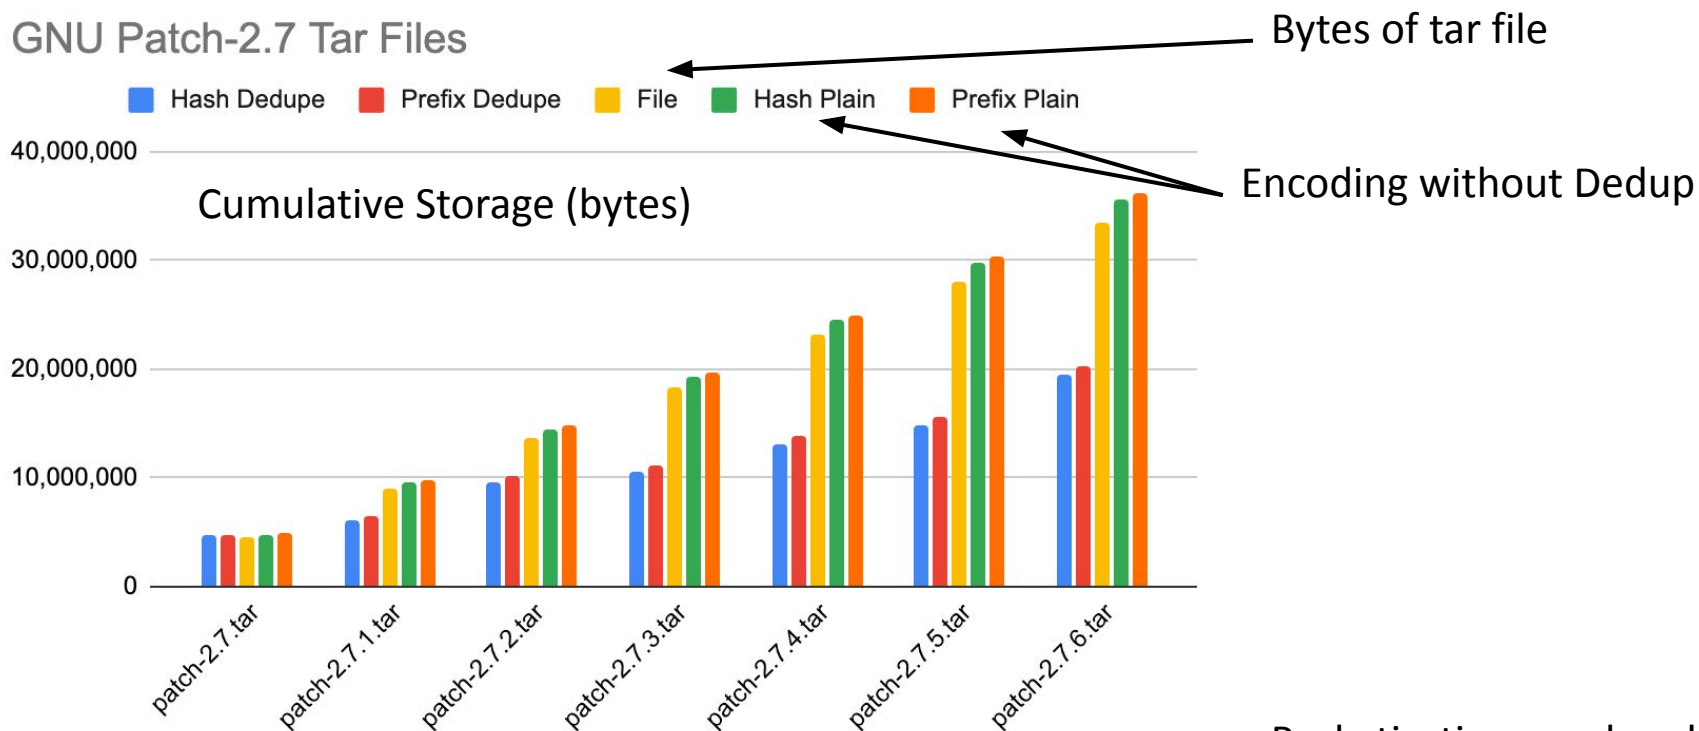
Naming

- The root manifest of each tar file has a name and signature.
- In HASHED
 - CCNx Nameless objects
- In PREFIX
 - /com/objectstore/gnu/d/ (+ hash)
 - /com/objectstore/gnu/m/ (+ hash)
 - differ by implicit hash.





Signed tar file manifest
points to top manifest
points to chunk manifests

GNU Patch-2.7 Tar Files




$$\% \text{ Savings} = \frac{(\text{ICN_size} - \text{file_size})}{\text{file_size}}$$

GNU emacs-29 tar files



GNU Bison-3.7 and Bison 3-8 Tar Files



Transition from
3.7 to 3.8

Savings takes
a hit

Experiment Results

- Encoding a series of releases within a project has significant savings.
- Savings
 - 30% to over 60% compared to file size.
 - Extra 5%-8% over ICN packetization.
 - About 2% - 3% Nameless over Prefix.

Unique Naming What-if

- tar root manifest
 - /objstore/gnu/patch/2.7.1/m
- tar internal manifests
 - /objstore/gnu/patch/2.7.1/m/<ManifestId>
- Chunk manifests
 - /objstore/gnu/patch/m/<?>
- Chunk data
 - /objstore/gnu/patch/d/<?>/<seqnum>

Conclusion

- The FLIC -07 draft should be ready for final review.
- We only have experience with CCNx and cefore forwarder.
- We suggested NDN encodings and would like to get NDN implementation experience.
- Upcoming IRTF drafts
 - Archives of multiple FLICs
 - Compression
 - Manifest diffs for updates